

# ADSA Data Analytics Project — Final Report

---

## Report index (Table of contents)

1. Project overview and objective
2. Step 1 — GitHub & code setup (expanded)
3. Step 2 — Error handling, logging & project structure
4. Step 3 — Data ingestion (implementation details)
5. Step 4 — Exploratory Data Analysis (EDA)
6. Step 5 — Data Transformation & Feature Engineering
7. Step 6 — Model Training (choices & evaluation)
8. Step 7 — Prediction / Inference pipeline
9. Reproducibility & how to run the pipeline
10. Results summary & discussion
11. Conclusion, limitations & future work
12. Appendices
  - A: Skeletal file structure
  - B: Sample scripts & snippets
  - C: Example config and requirements

---

## 1. Project overview and objective

**Goal.** Build a reproducible ML pipeline that ingests movie (or given) metadata, performs EDA, cleans and transforms the data, trains and validates models, and exposes a prediction pipeline. Deliverables: working code (src/), documented GitHub repo, reproducible environment, and a short report containing results and next steps.

**Data.** Describe dataset(s) used (columns, number of rows, source). If multiple files (CSV/BigQuery), list them here.

**Success criteria.** e.g., target metric (RMSE, MAE, R2, Accuracy), baseline performance, and production readiness requirements (unit tests, logging, packaging).

---

## 2. Step 1 — GitHub & code setup (expanded)

### 2.1 Environment creation (conda + venv)

Use an isolated environment to ensure reproducibility. Create a conda environment in the project folder and activate it.

Typical commands to run:

- `conda create -p ./venv python==3.8 -y`
- `conda activate ./venv`
- Install dependencies from `requirements.txt` after creating it: `pip install -r requirements.txt`

Note: Use explicit pinned versions in `requirements.txt` for reproducibility.

## 2.2 Git initialization & remote

Initialize a repository, create README and .gitignore, add files, commit, set the main branch and add a remote.

Typical git commands to run:

- `git init`
- `git add README.md .gitignore`
- `git commit -m "first commit"`
- `git branch -M main`
- `git remote add origin https://github.com/Manoj-Sh-AI/ADSA_Data_Analytics-1_ml_pipeline.git`
- `git push -u origin main`

## 2.3 Project packaging (setup.py)

If you plan to install the project as a package during development, create a `src/` package and a minimal `setup.py`. Use `pip install -e .` to install locally during development.

## 2.4 Requirements

Keep `requirements.txt` updated and pin versions. Example packages to include: pandas, numpy, scikit-learn, xgboost, catboost, joblib, pyyaml, pytest.

---

# 3. Step 2 — Error handling, logging & project structure

## 3.1 Project structure (high-level)

(See Appendix A for full skeletal tree.)

Key folders and responsibilities:

- `src/components/` : modular units (data\_ingestion, data\_transformation, model\_trainer)
- `src/pipeline/` : orchestrates end-to-end flow (train\_pipeline, predict\_pipeline)
- `src/config/` : YAML/JSON with paths, hyperparameters, and constants
- `src/logs/` : runtime logs
- `tests/` : unit & integration tests

## 3.2 Logging

Use Python `logging` configured to write to both console and a file in a `logs/` folder. Provide a helper (e.g., `get_logger`) that sets up formatters, stream handler and file handler. Call the logger at the top of each module to standardize messages and timestamps.

## 3.3 Exception handling

Create a custom exception class to attach context and preserve tracebacks. Wrap high-level pipeline functions with try/except blocks, log the exception with stack trace (using `logger.exception`), and re-raise or handle with meaningful messages.

---

# 4. Step 3 — Data Ingestion

## 4.1 Goals

- Read raw data from CSV / BigQuery / local storage
- Validate schema (expected columns)
- Save raw snapshot (versioned) and a cleaned copy for transformations

## 4.2 Data ingestion API (recommended behaviour)

Design data ingestion to:

- Accept configuration (paths for raw, train and test)
- Read the raw dataset
- Perform a deterministic train/test split (e.g., 80/20, with a fixed seed)
- Persist the train and test sets to `data/processed/`
- Log operations and raise meaningful exceptions on failure

## 4.3 Validation checks

- Column presence
- Missing value proportions per column
- Duplicate rows
- Data types and ranges (e.g., dates, numeric ranges)

---

## 5. Step 4 — Exploratory Data Analysis (EDA)

### 5.1 Objectives

- Understand distributions, missingness, correlation with target
- Detect outliers and data quality issues
- Identify candidate feature transformations

### 5.2 Recommended EDA checklist

- Summary statistics (describe and info)
- Missing value heatmap and percent-missing table
- Univariate plots for numeric features (histograms / boxplots)
- Countplots for categorical variables
- Correlation matrix / heatmap for numeric relationships
- Target vs feature plots (scatter, violin, bar) for top predictors
- Time-series checks if data is temporal

### 5.3 Tools

Pandas and Matplotlib (or Seaborn for convenience) are recommended. Conduct EDA interactively in Jupyter notebooks and save relevant visuals to `reports/figures/`.

---

## 6. Step 5 — Data Transformation & Feature Engineering

### 6.1 Goals

- Convert raw features into model-ready features
- Handle missing data, encode categoricals, scale numeric features
- Persist transformers (scalers, encoders) with joblib for inference

### 6.2 Typical steps

1. **Imputation**: numeric (mean/median/KNN), categorical (mode or new category "Missing")
2. **Categorical encoding**: One-Hot for low-cardinality, Target/Ordinal or Count encoding for high-cardinality
3. **Scaling**: Standard scaling for models sensitive to scale, MinMax for neural networks
4. **Feature creation**: interaction terms, datetime decompositions, aggregated features
5. **Dimensionality reduction** (if needed): PCA, feature selection via tree-based importance

## 6.3 Implementation notes

- Build a preprocessing pipeline object that encapsulates imputation, encoding and scaling and fit it on the training data only.
  - Persist the fitted preprocessing object to `artifacts/` so it can be reused by the prediction pipeline.
- 

# 7. Step 6 — Model Training (Why I chose these models and how I achieved the accuracy)

## 7.1 Modeling strategy

- Try multiple models: Linear Regression (baseline), RandomForest, XGBoost, CatBoost, and a simple neural net if needed.
- Use cross-validation (KFold, or TimeSeriesSplit if temporal) and compare with metrics (RMSE, MAE, R2). Keep a validation set for final selection.

## 7.2 Why these models?

- **Linear Regression**: interpretable baseline, fast.
- **RandomForest**: strong baseline for tabular data, robust to outliers and feature scaling.
- **XGBoost/CatBoost**: state-of-the-art for many tabular problems; often give best performance with moderate tuning.

## 7.3 Hyperparameter tuning

Use randomized or grid search for hyperparameter tuning; for faster, use Bayesian optimizers (Optuna). Keep `n_jobs=-1` where possible and use sensible search spaces.

## 7.4 Evaluation

Report train/validation/test metrics and show learning curves. Record metrics in a `reports/` file and compare several candidate models in a table.

## 7.5 Persisting the model

Persist the trained model artifact to `artifacts/` (e.g., `model.joblib`) for later inference.

---

## 8. Step 7 — Prediction / Inference pipeline

### 8.1 Design

- Provide a single entry point (predict pipeline) that accepts raw input (CSV or JSON), applies the saved preprocessing object, loads the model artifact, and outputs predictions (CSV/JSON).
- Validate input schema and produce outputs with predictable column names (e.g., `prediction`).

### 8.2 Inference behaviour

- Load preprocessor and model artifacts from `artifacts/`.
- Transform the raw input using the preprocessor and call `model.predict()`.
- Return or persist predictions alongside input identifiers.

### 8.3 Deployment considerations

- For a demo: expose via a lightweight API (Flask/FastAPI).
- For production: containerize, add health checks, secure secrets, and add monitoring.

---

## 9. Reproducibility & How to run the pipeline

### 9.1 Quick start (development)

Typical steps to run locally:

1. Clone the repo.
2. Create the environment and install requirements.
3. `conda create -p ./venv python==3.8 -y`
4. `conda activate ./venv`
5. `pip install -r requirements.txt`
6. Run the training pipeline (script entry point) and the prediction script with input and output paths.  
Ensure `config/` contains correct paths and seeds.

---

## 10. Results summary & discussion

Include here a short table with final metrics for the chosen model and the baseline. Also include: feature importances, top 5 features, and saved visuals (learning curves, residual plots) placed in `reports/figures/`.

Example columns to report: Model, Dataset, RMSE, MAE, R2.

## 11. Conclusion, limitations & future work

- **Conclusion:** Which model you selected and why, summary of results.
  - **Limitations:** data quantity/quality, label noise, potential leakage, or limited features.
  - **Future work:** more feature engineering, ensembling, better cross-validation, model explainability (SHAP), productionizing with CI/CD.
- 

## 12. File Structure

### Appendix A — Skeletal file structure

```
ADSA_Data_Analytics-1_ml_pipeline/
├── .gitignore
├── README.md
├── requirements.txt
├── setup.py
├── config/
│   └── config.yaml
├── data/
│   ├── raw/
│   ├── processed/
│   └── external/
├── src/
│   ├── __init__.py
│   ├── logger.py
│   ├── exception.py
│   ├── utils.py
│   ├── components/
│   │   ├── __init__.py
│   │   ├── data_ingestion.py
│   │   ├── data_transformation.py
│   │   └── model_trainer.py
│   ├── pipeline/
│   │   ├── __init__.py
│   │   ├── train_pipeline.py
│   │   └── predict_pipeline.py
│   └── config/
│       └── params.yaml
├── artifacts/
│   ├── preprocessor.joblib
│   └── model.joblib
└── reports/
```

```
| |─ figures/
| |─ report.pdf
|─ tests/
|   └─ test_data_ingestion.py
```

## Appendix B — Sample scripts & snippets

- `train_pipeline` should: load config → run data\_ingestion → run data\_transformation → run model\_trainer → save artifacts and metrics.
  - `predict_pipeline` should: parse args → validate input → load artifacts → run predict → save results.
-